

Applied Data Science with R Capstone project

Joshua Agbroko

January 2024

Outline



- Executive Summary
- Introduction
- Methodology
- Results
- Conclusion
- Appendix

Executive Summary



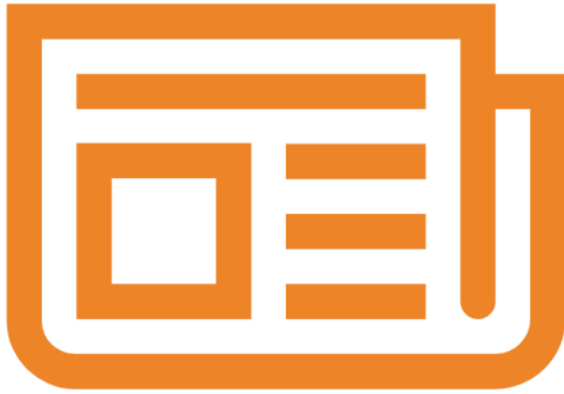
- This research tries to analyze how weather would affect bike sharing demand in urban areas. The dataset used in this analysis is from Seoul Bike Sharing System, Open Weather API, World Cities, and Bike System in the World
- We developed a predictive model to forecast bike demand under varying weather conditions by leveraging:
 - Historical weather
 - Bike Rental records from multiple cities
- By integrating weather parameters such as temperature, humidity, wind speed, and precipitation, our model provide insights into how weather influences bike usage patterns.

Introduction



- Bike sharing systems are a type of bicycle rental service in which the procedure of obtaining a membership, renting a bike, and returning the bike is all done through a network of kiosks located around a city.
- People can rent a bike from one location and return it to a different location on an as-needed basis using these systems
- In urban environments, bike-sharing systems play a pivotal role in promoting sustainable transportation and addressing congestion challenges.
- However, optimizing bike-sharing services requires a deep understanding of factors influencing demand, particularly the impact of weather conditions.
- Challenges faced include;
 - Dynamic nature of demands
 - Weather sensitivity
 - Optimization needs

Methodology

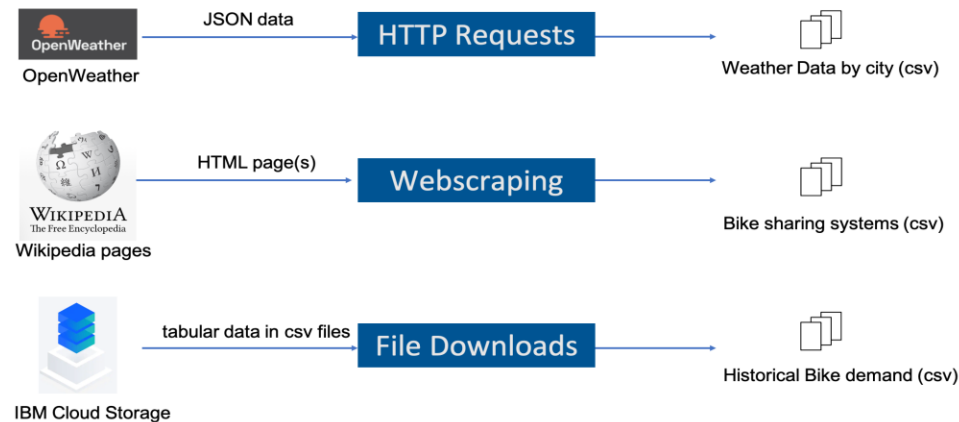


- Perform data collection
- Perform data wrangling
- Perform Exploratory Data Analysis (EDA) using SQL and visualization
- Perform predictive analysis using regression models
 - How to build the baseline model
 - How to improve the baseline model
- Build a R Shiny dashboard app

Methodology

Data collection

- Weather Data was Retrieved from OpenWeatherAPI Using OpenWeather REST API to request weather data.
- Web scraping information about global bike-sharing systems from Wikipedia web pages, and downloading and aggregating tabular data from cloud storage.



Data wrangling

- Many related inconsistencies and noises like inconsistent data formats, extracted fields containing HTML tags, reference links were eliminated using regular expressions.
- Detected and handled missing values that may affect the model's predictive ability.
- Categorical variables were converted into indicator variables
- Numeric columns were normalized to transfer them all into a similar range
- Add the screenshots of data wrangling code cell and output for regular expressions, missing values handling, generating indicator columns to the Appendix section for peer-review

EDA with SQL

- Performed SQL queries using RSQLite to :
 - Find valuable statistics such as total bike-sharing records or operational hours,
 - Filter data based on city or date,
 - Find patterns such as bike-sharing seasonality or similarities among bike-sharing systems across the world.

EDA with data visualization

- Understanding the distributions of the data using histograms
- Finding correlations between important features using scatterplots
- Spotting outliers and irregular behavior in your features using box plots.

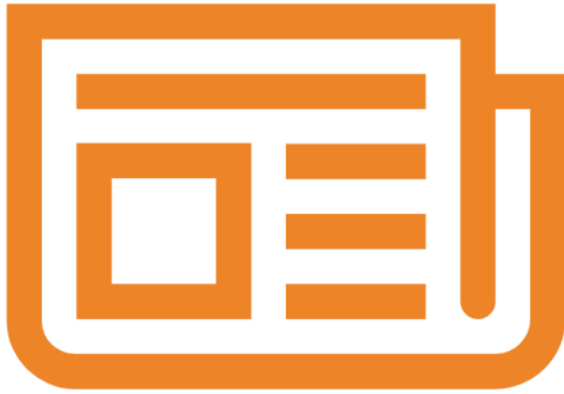
Predictive analysis

- Embarked on building a predictive model to forecast bike-sharing demand, a crucial task for urban transportation planning and resource allocation. Began by collecting extensive weather data from the OpenWeatherAPI and integrating it with historical bike usage records from various cities. Leveraging a diverse set of machine learning algorithms including linear regression, random forest, gradient boosting, we constructed initial models to predict bike demand. These models were evaluated using key metrics such as Root Mean Squared Error (RMSE) and R-squared to gauge their accuracy and performance.

Build a R Shiny dashboard

- Integrated interactive leaflet maps displaying cities' bike-sharing demand predictions with color-coded markers indicating predicted demand levels.
- Implemented dropdown menu functionality allowing users to select specific cities for detailed analysis.
- Incorporated line plots depicting temperature forecasts for the next five days, aiding in understanding weather patterns' influence on bike demand.
- Developed line plots illustrating bike-sharing demand predictions over time, facilitating trend analysis and forecasting assessment.

Results



- Exploratory data analysis results
- Predictive analysis results
- A dashboard demo in screenshots

EDA with SQL


Busiest bike rental times

- DATE HOUR total_rentals

19/06/2018 18 3556




- This query result indicates that June 19, 2018 6:00 PM, had the most bike rentals recorded in the dataset.

Hourly popularity and temperature by seasons

Result Grid   Filter Rows: <input type="text"/> Export:  Wrap Cell Content: 			
	SEASONS	avg_hourly_temperature	avg_bike_rentals_per_hour
▶	Summer	26.587710514337775	1034.0733695652175
	Autumn	13.82157976251933	924.1104801239029
	Spring	13.021685185185188	746.2541666666667
	Winter	-2.5404629629629607	225.5412037037037

- This analysis highlights the dynamic relationship between seasonal weather patterns, temperature, and bike rental activity, underscoring the need for data-driven decision-making in the bike-sharing industry

Rental Seasonality

• Result Grid			Filter Rows: <input type="text"/>	Export: 	Wrap Cell Content: 
	SEASONS	avg_hourly_bike_count	min_hourly_bike_count	max_hourly_bike_count	std_dev_hourly_bike_count
▶	Winter	225.5412037037037	3	937	150.3374234675032
	Spring	746.2541666666667	2	3251	618.5247350747628
	Summer	1034.0733695652175	9	3556	690.0884362677946
	Autumn	924.1104801239029	2	3298	617.3884505672105

These observations provide insights into seasonal bike usage patterns, which can be valuable for bike-sharing system optimization. The observations underline the fact that far less bikes are rented during Winter.

Weather Seasonality

SEASONS	avg_temperature	avg_humidity	avg_wind_speed	avg_visibility	avg_dew_point_temperature	avg_solar_radiation	avg_rainfall	avg_snowfall	avg_bike_count
Summer	26.587710514337775	64.9814311594203	1.6094202898550731	1501.7454710144928	18.75013586956522	0.7612545289855057	0.2534873188405798	0	1034.0733695652175
Autumn	13.82157976251933	59.044914816726894	1.4921011874032002	1558.1744966442952	5.150593701600414	0.5227826535880221	0.11765616933402169	0.06350025813113062	924.1104801239029
Spring	13.021685185185188	58.75833333333333	1.8577777777777793	1240.911574074074	4.091388888888886	0.6803009259259252	0.1869444444444444	0	746.2541666666667
Winter	-2.5404629629629607	49.74490740740741	1.9226851851851832	1445.987037037037	-12.41666666666665	0.2981805555555558	0.032824074074075	0.2474999999999997	225.5412037037037

Based on these observations, we can deduce that weather conditions, particularly temperature and solar radiation, may influence bike-sharing usage, with higher usage observed during warmer and sunnier seasons. Additionally, factors such as humidity, visibility, and precipitation may also impact bike-sharing patterns, albeit to a lesser extent.

Bike-sharing info in Seoul

CITY	COUNTRY	LAT	LO	POPULATION	total_bikes_available
1 Seoul	Korea, South	37.5833	127	21794000	20000

- This query result provides insight into the population size and availability of bicycles for shared use in Seoul.

Cities similar to Seoul

```
print(conn)

CITY      COUNTRY  LAT    LNG POPULATION total_bikes
1 Beijing    China  39.9050 116.3914  19433000    16000
2 Ningbo     China  29.8750 121.5492   7639000    15000
3 Shanghai   China  31.1667 121.4667  22120000    19165
4 Weifang    China  36.7167 119.1000   9373000    20000
5 Xi'an      China  34.2667 108.9000   7135000    20000
6 Zhuzhou    China  27.8407 113.1469   3855609    20000
7 Seoul      Korea, South 37.5833 127.0000  21794000    20000

close(conn)
```

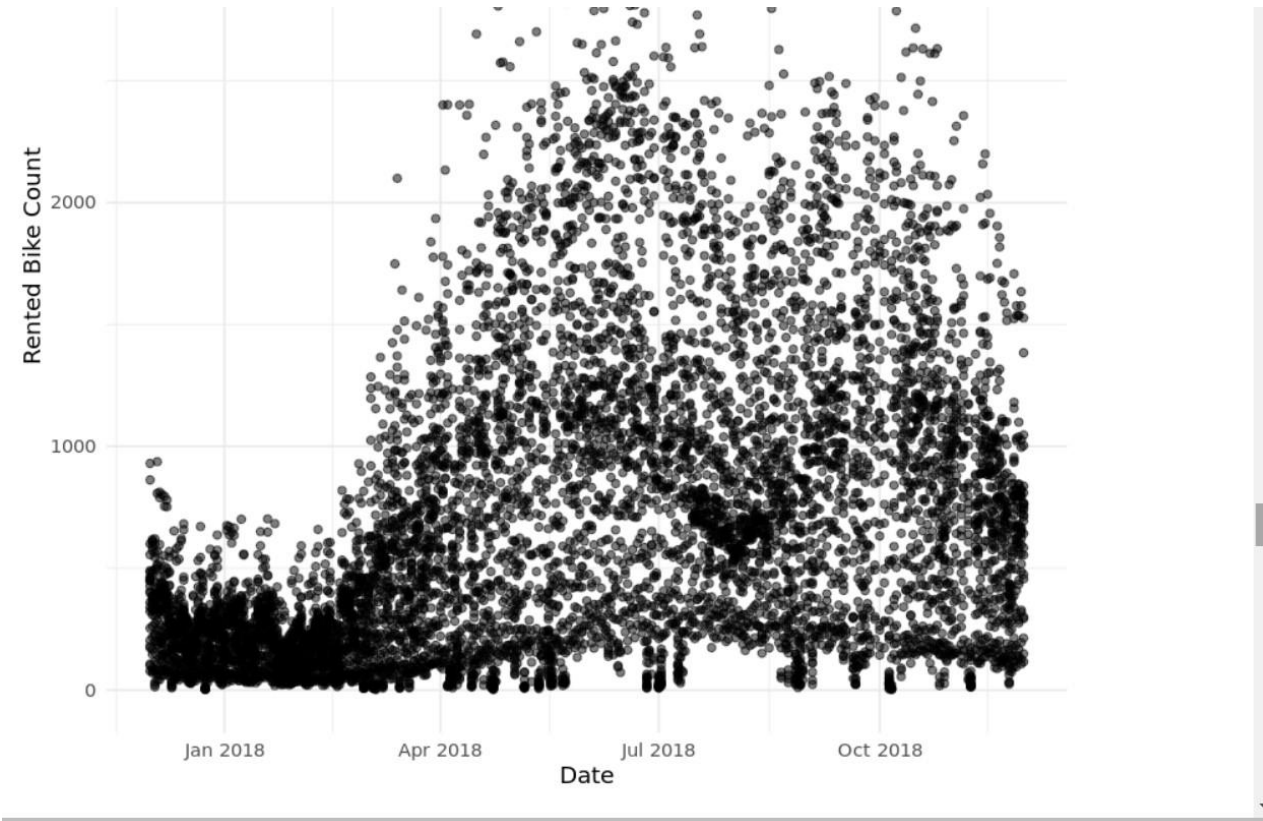
We can infer that these cities, especially Beijing and Shanghai likely have active bike-sharing systems or initiatives due to their large populations and significant numbers of available bicycles just like Seoul.

EDA with Visualization

Bike rental vs. Date

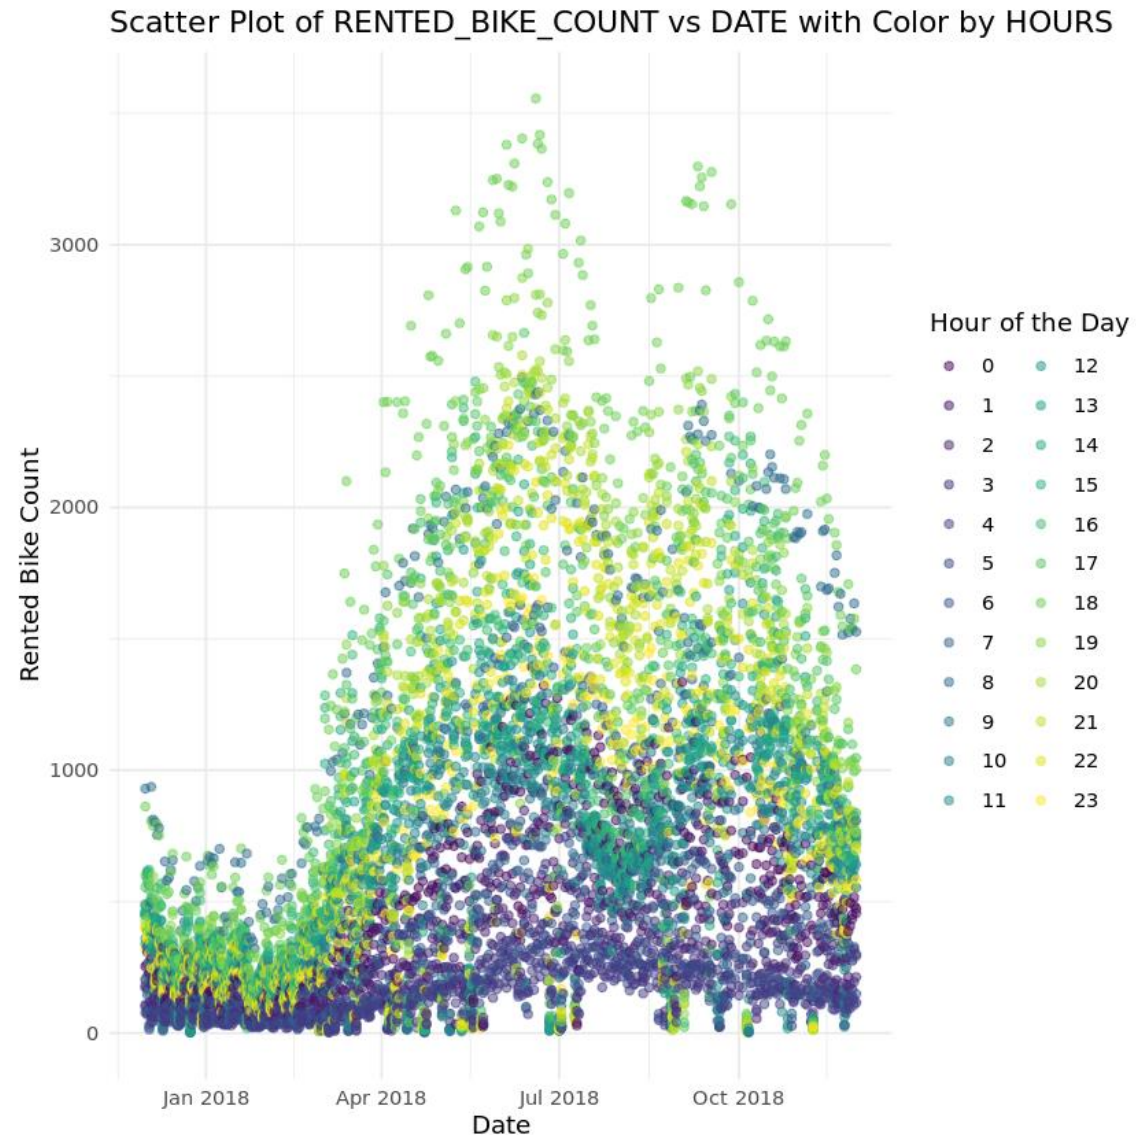
Scatter plot
of RENTED_BIKE_COUNT vs. DATE

Noticeably few bike rentals during the winter months. Sharp increase during the following Spring months which has the highest rentals. Summer and Autumn also have high number of bike rentals.



Bike rental vs. Datetime

Afternoon to late evening has the highest number of bike rentals with nighttime as expected, having the lowest numbers

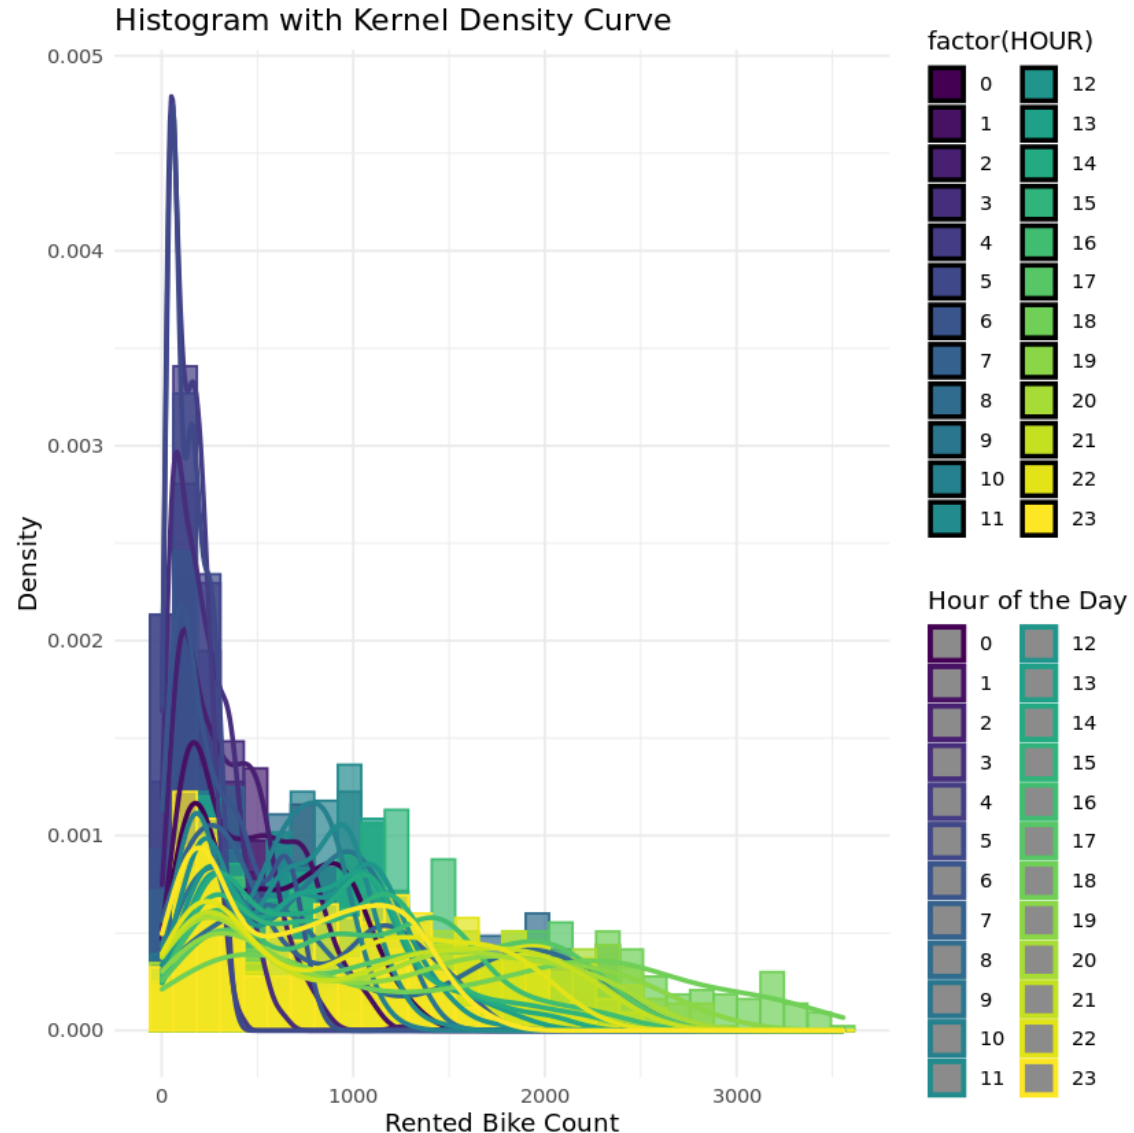


Bike rental histogram

We can see from the histogram that most of the time there are relatively few bikes rented. Indeed, the 'mode', or most frequent amount of bikes rented, is about 250.

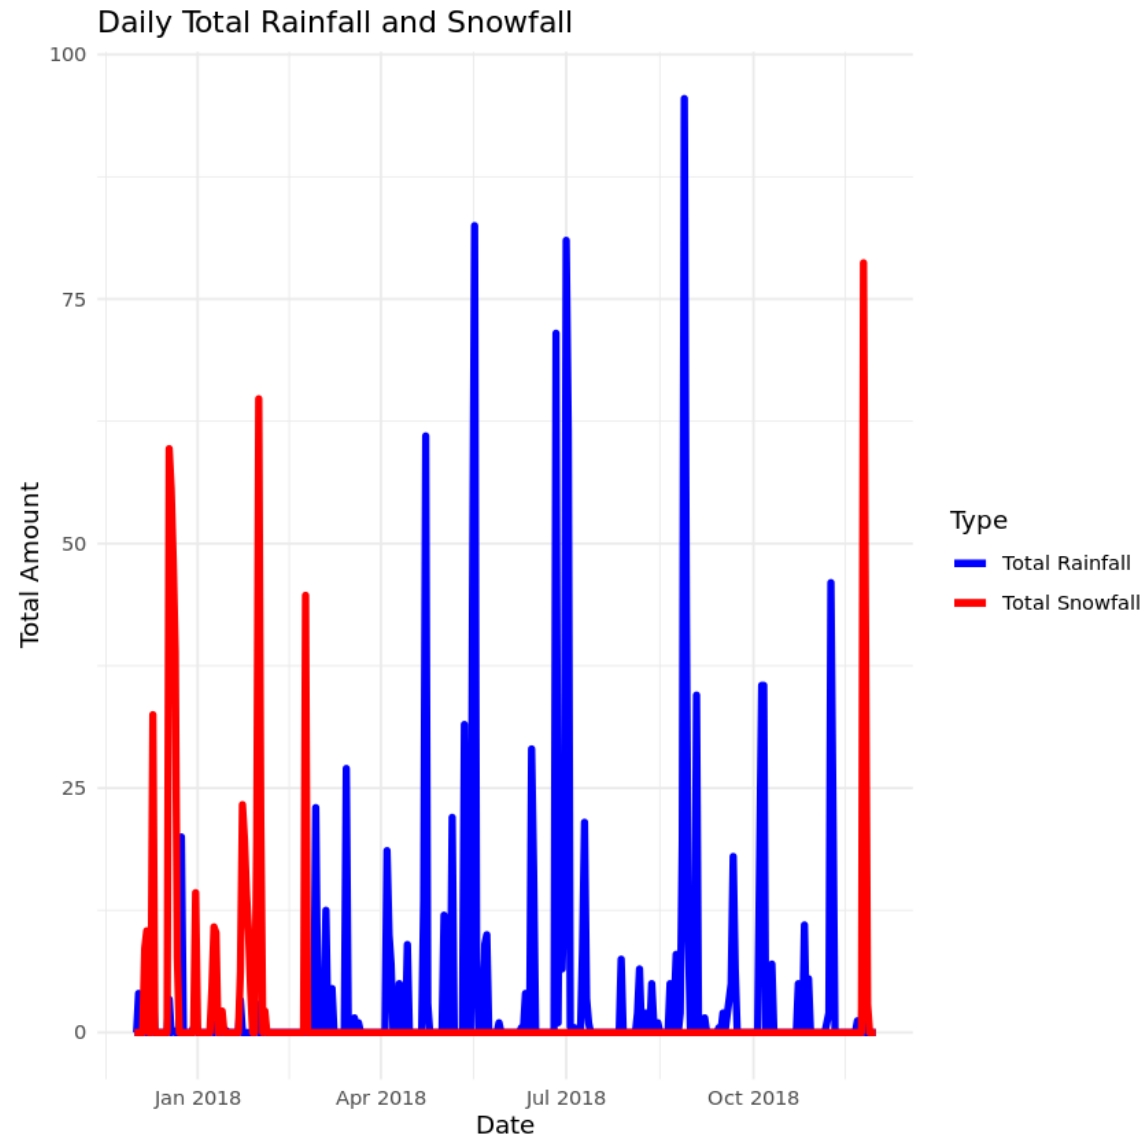
Judging by the 'bumps' at about 700, 900, and 1900, and 3200 bikes, it looks like there may be other modes hiding within subgroups of the data.

Interestingly, judging from the tail of the distribution, on rare occasions there are many more bikes rented out than usual.



Daily total rainfall and snowfall

A barchart calculating the daily total rainfall and snowfall

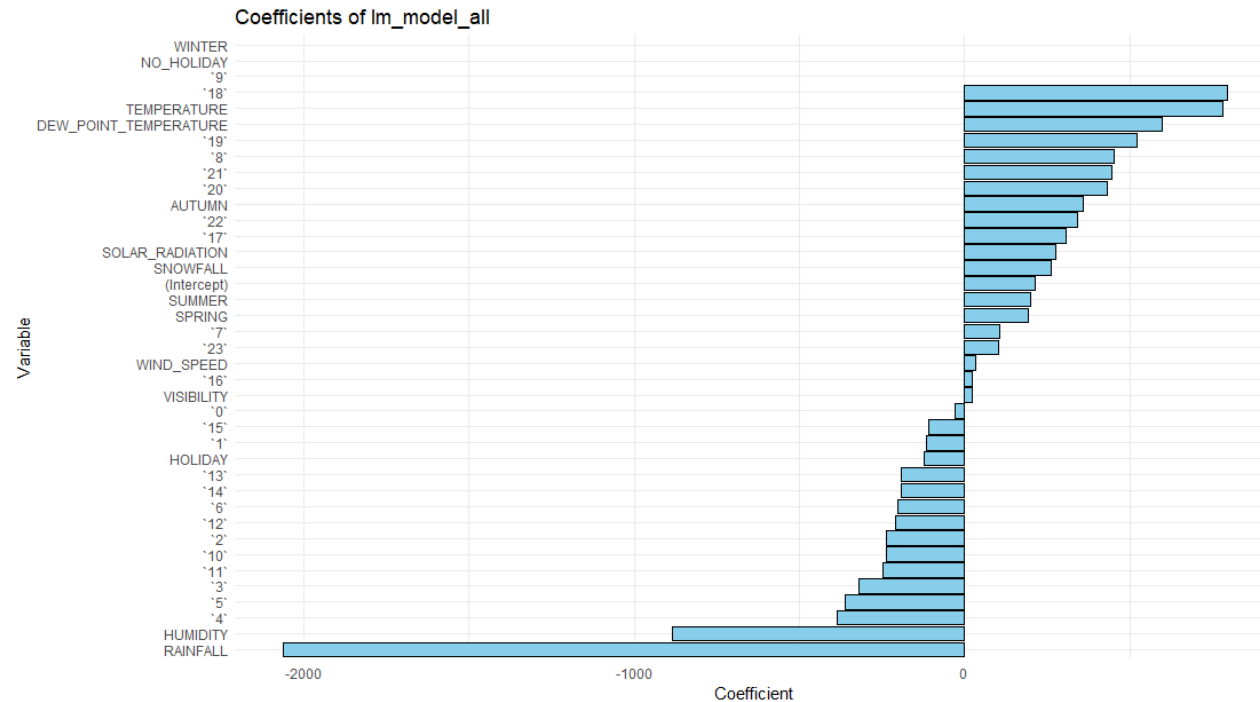


Predictive analysis

Ranked coefficients

From the coefficients of the model, we can find that Most of the highest predictors are weather variables such as Rainfall, Humidity, Temperature, and Dew Point Temperature. It means weather condition could influence people decision to rent bikes.

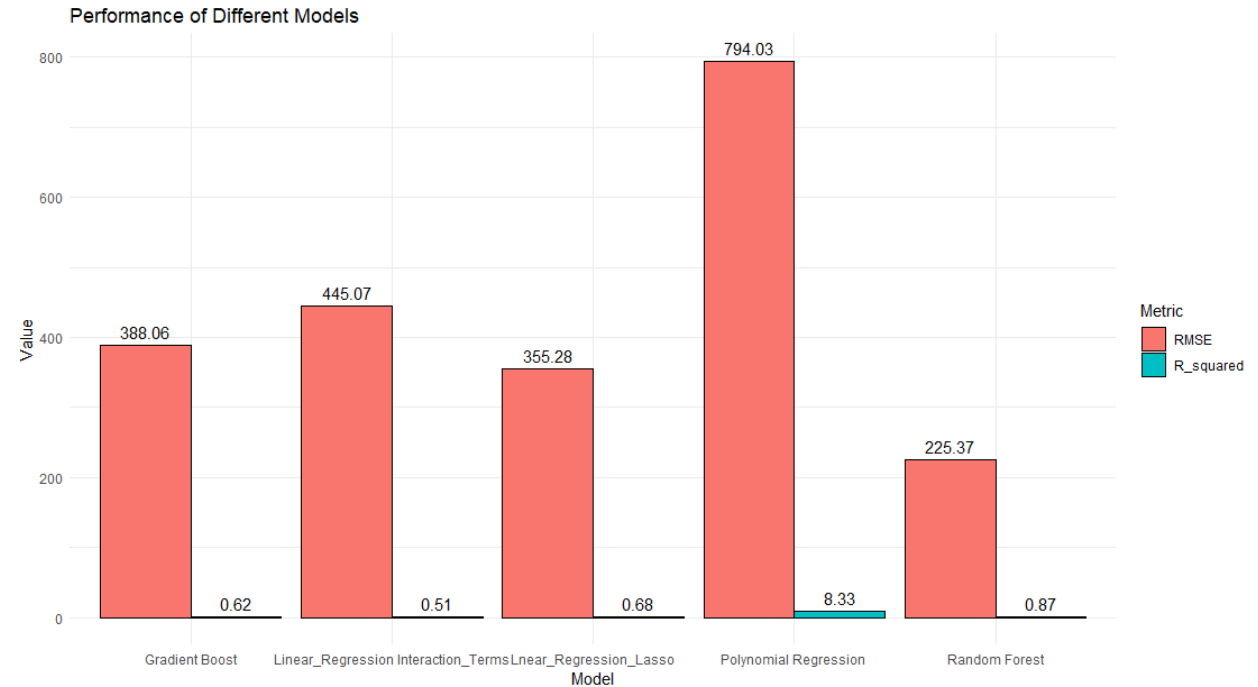
Another significant indicator is most evening time is highly correlated with higher number of bike rents



Model evaluation

► Here is the result of RMSE and RSquared of each models created for the estimation

► Visualized are the refined models' RMSE and R-squared using grouped bar chart



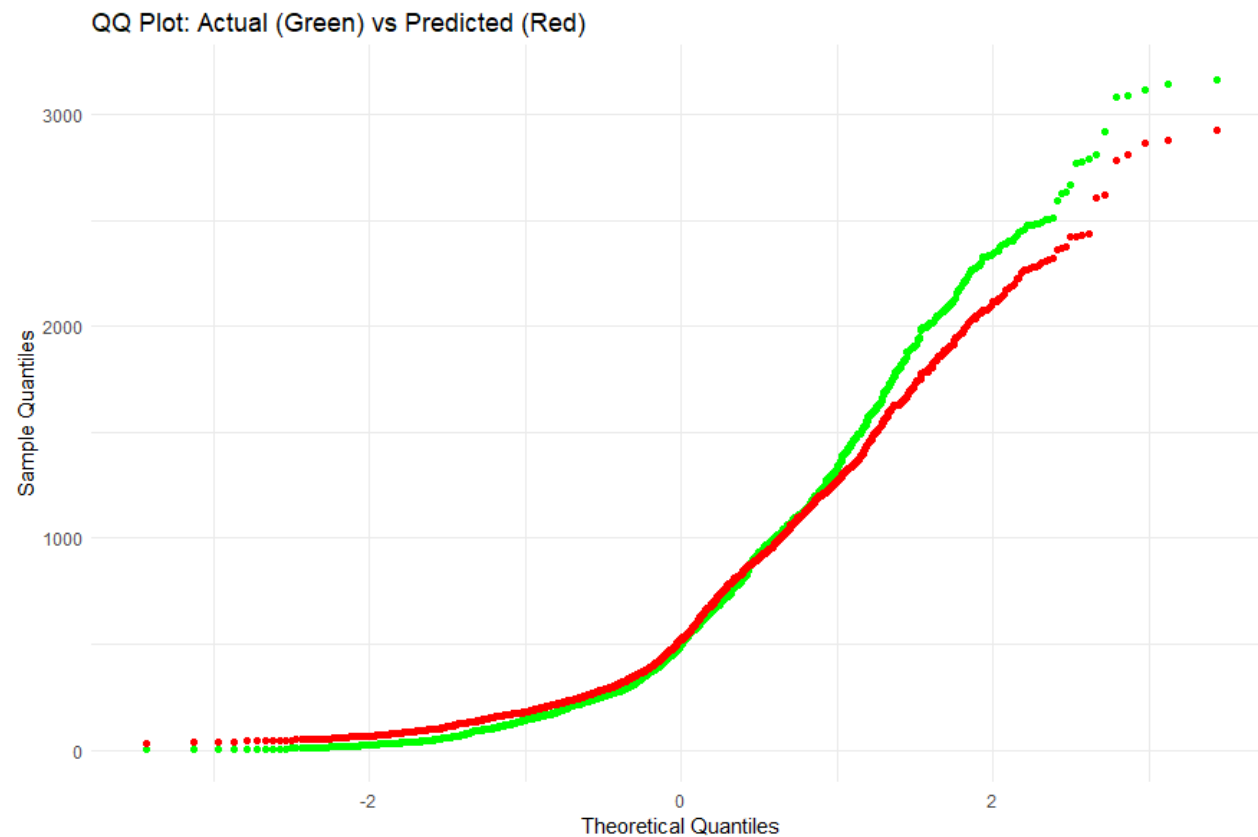
Find the best performing model

```
> rmse_rf  
[1] 225.3694  
> rsquared_rf  
[1] 0.873195  
> |
```

- The predictors <- model.matrix(RENTED_BIKE_COUNT ~ TEMPERATURE + HUMIDITY + WIND_SPEED + VISIBILITY + DEW_POINT_TEMPERATURE + SOLAR_RADIATION + RAINFALL + SNOWFALL, data = train_data1)[, -1]
- response <- train_data\$RENTED_BIKE_COUNT
- rf_model <- randomForest(predictors, response)

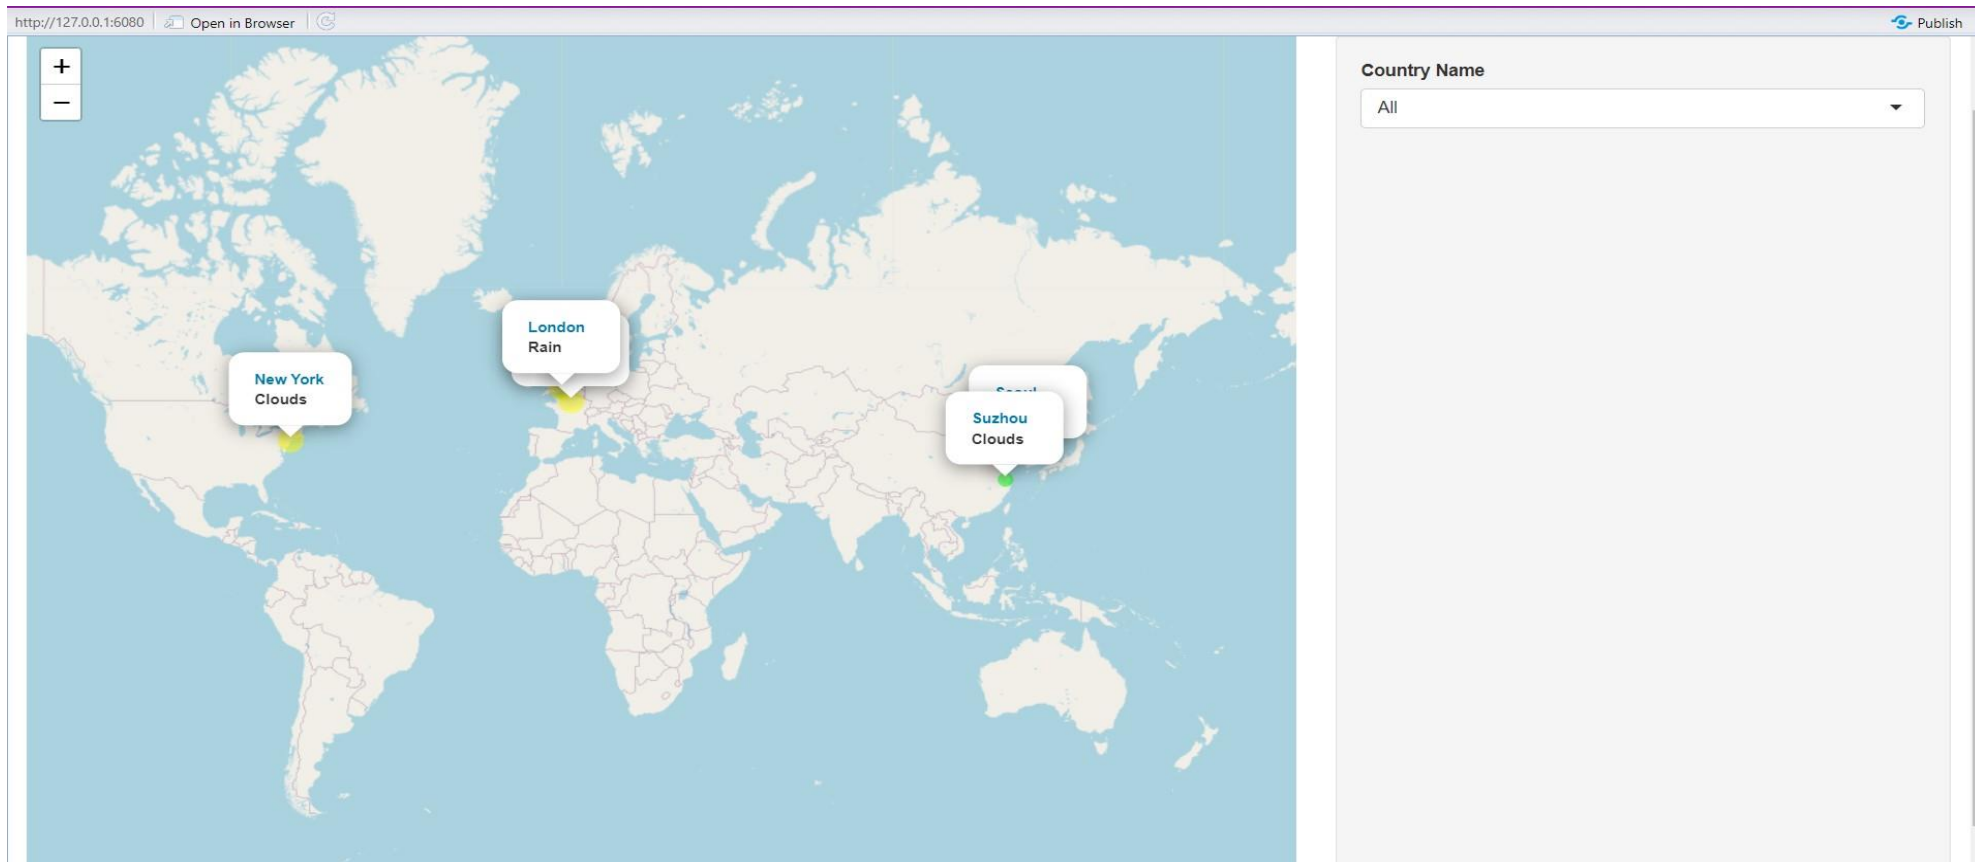
Q-Q plot of the best model

the Q-Q plot of the best model's test results
vs the truths



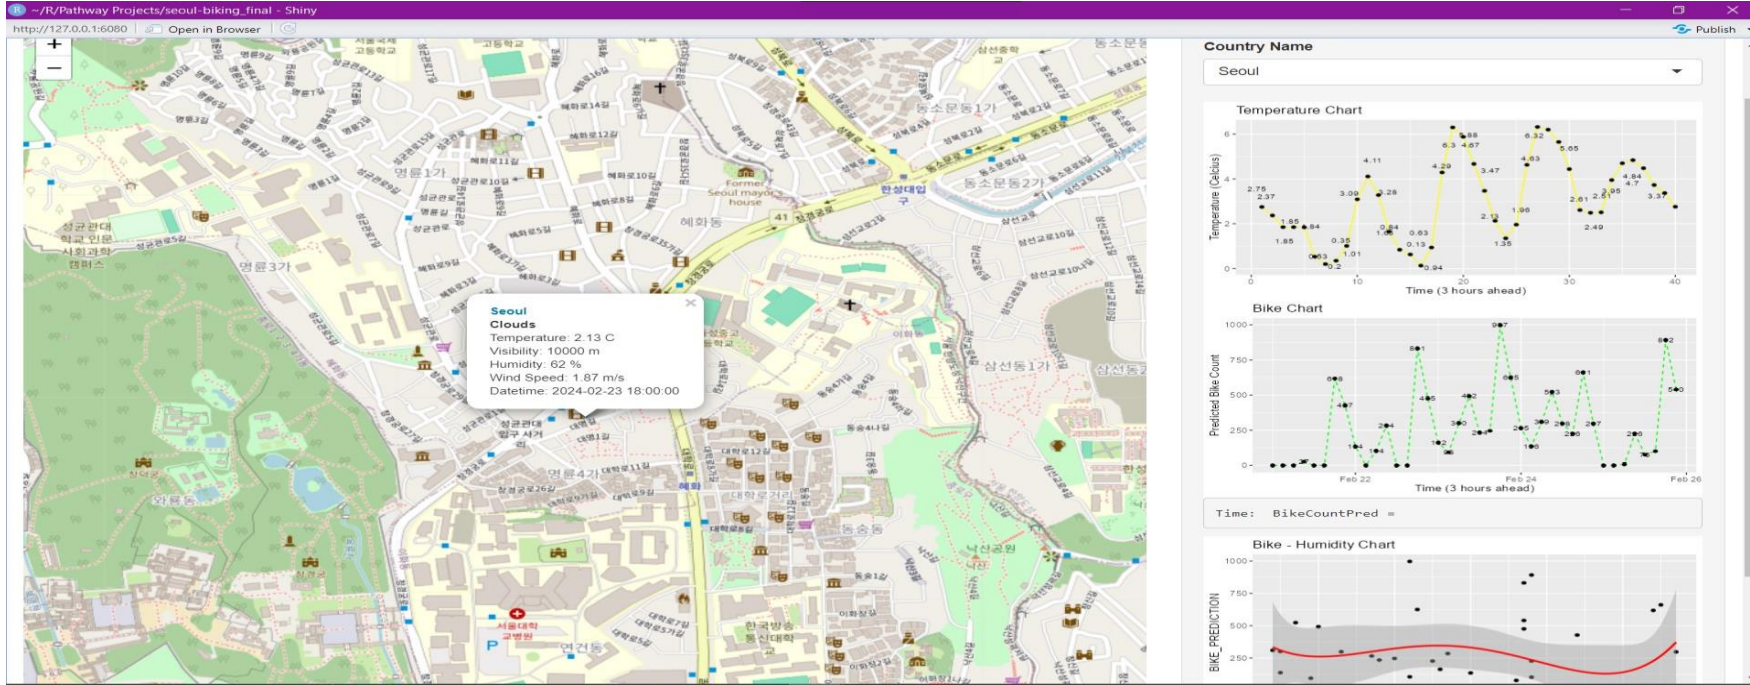
Dashboard

Bike-sharing Demand Prediction Dashboard Overview



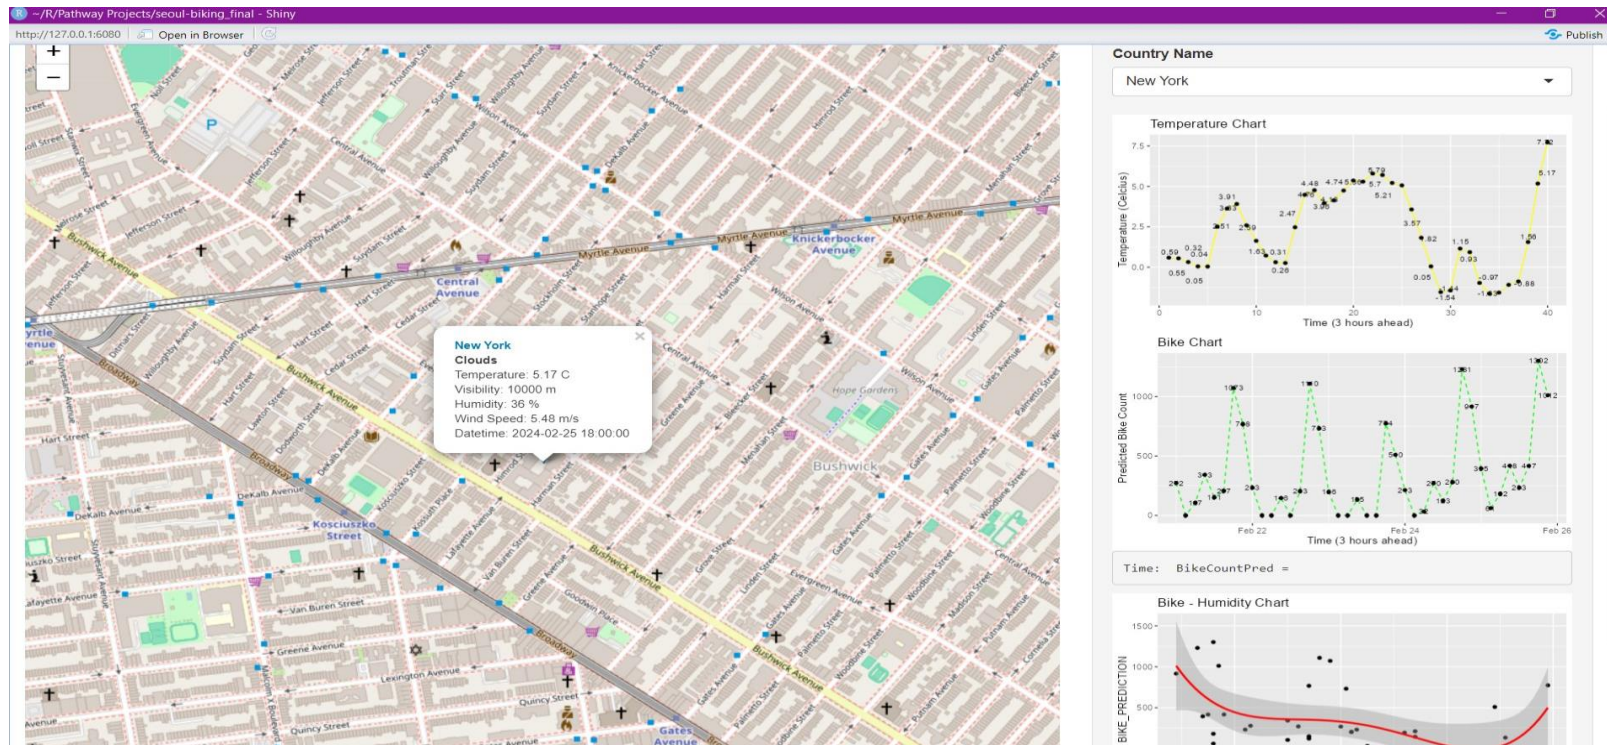
Main default page of Shiny app showing general overview of all cities from which we can select a specific city to show more details about each of its bike sharing demand information.

Selected City : SEOUL



Specific city, Seoul, selected to show its temperature, bike prediction count and relative humidity at specific time intervals. Closer city view of Seoul also shown on map.

Selected City : New York



Another city, this time New York, is selected on the shiny app showing this cities specific temperature, humidity and bike prediction count data with a closer view of the city


CONCLUSION



- The observed patterns in bike-sharing demand highlight the significant influence of seasons on rental behavior. Rental activity peaks during the summer and autumn seasons, coinciding with warmer temperatures and favorable weather conditions, while experiencing a decline during winter and early spring. This seasonal variation underscores the importance of considering environmental factors and seasonal trends when designing bike-sharing services and infrastructure. Understanding how seasonality affects rental behavior can inform strategic decisions regarding service expansion, promotional campaigns, and resource allocation to meet fluctuating demand throughout the year.
- The insights gleaned from the model should enable stakeholders to make informed decisions, such as adjusting bike inventory levels based on weather forecasts, time of day patterns, and seasonality trends, thereby enhancing the overall user experience of bike-sharing systems.
- Moving forward, there are several avenues for further exploration and improvement. Enhancing the model's predictive accuracy by incorporating additional data sources, such as demographic information, traffic patterns, and special events, could provide deeper insights into bike-sharing dynamics.

APPENDIX

- Included are all relevant assets like R code snippets, SQL queries, charts,.



Data Collection (Weather API)

TODO: Get the root HTML node

```
[2]: url <- "https://en.wikipedia.org/wiki/List_of_bicycle-sharing_systems"  
      # Get the root HTML node by calling the `read_html()` method with URL
```

```
[3]: webpage <- read_html(url)  
  
      table_nodes <- html_nodes(webpage, "table")
```

```
[5]: if (length(table_nodes) >= 1) {  
      first_table_data <- html_table(table_nodes[[1]], fill = TRUE)  
  
      # Print the extracted data from the first table as a data frame  
      df_first_table <- as.data.frame(first_table_data)  
      print(df_first_table)  
    } else {  
      cat("No tables found on the webpage.\n")  
    }
```

	Country
1	Albania
2	Argentina
3	Argentina
4	Argentina
5	Argentina

Data Collection (Weather API)

Summarize the bike sharing system data frame

```
[6]: # Summarize the dataframe  
summary(df_first_table)
```

Country	City	Name	System
Length:564	Length:564	Length:564	Length:564
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
Operator	Launched	Discontinued	Stations
Length:564	Length:564	Length:564	Length:564
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
Bicycles	Daily ridership		
Length:564	Length:564		
Class :character	Class :character		
Mode :character	Mode :character		

Export the data frame as a csv file called `raw_bike_sharing_systems.csv`

```
[7]: # Export the dataframe into a csv file  
write.csv(df_first_table, file = "raw_bike_sharing_systems.csv", row.names = FALSE)  
  
# Print a message indicating the successful export  
cat("Data frame exported to raw_bike_sharing_systems.csv\n")  
  
Data frame exported to raw_bike_sharing_systems.csv
```

Data Collection (Webscrapping)

```
In [5]: # Converting the bike-sharing system table into a dataframe
table_nodes <- html_nodes(webpage, "table")

# Extracting information from the first table
if (length(table_nodes) >= 1) {
  first_table_data <- html_table(table_nodes[[1]], fill = TRUE)

  # Printing the extracted data from the first table as a data frame
  df_first_table <- as.data.frame(first_table_data)
  print(df_first_table)
} else {
  cat("No tables found on the webpage.\n")
}
```

```
1          Country
2      Albania
3      Argentina
4      Argentina
5      Argentina
6      Australia
7      Australia
8      Australia
9      Australia
10     Australia
11     Australia
12         Austria
13         Austria
14         Austria
15         Austria
16         Austria
17         Austria
18     Bangladesh
19         Belgium
20         Belgium
21         Belgium
```

```
6]: # Summarize the dataframe
summary(df_first_table)

Country      City      Name      System
Length:560   Length:560   Length:560   Length:560
Class :character  Class :character  Class :character  Class :character
Mode :character  Mode :character  Mode :character  Mode :character
Operator      Launched      Discontinued      Stations
Length:560   Length:560   Length:560   Length:560
Class :character  Class :character  Class :character  Class :character
Mode :character  Mode :character  Mode :character  Mode :character
Bicycles      Daily ridership
Length:560   Length:560
Class :character  Class :character
Mode :character  Mode :character
```

Export the data frame as a csv file called `raw_bike_sharing_systems.csv`

```
7]: # Exporting the dataframe into a csv file
write.csv(df_first_table, file = "raw_bike_sharing_systems.csv", row.names = FALSE)

# Print a message indicating the successful export
cat("Data frame exported to raw_bike_sharing_systems.csv\n")

Data frame exported to raw_bike_sharing_systems.csv
```

Data Wrangling (Regex)

To improve dataset readability by both human and computer systems, we first need to standardize the column names of the datasets above using the following naming convention:

- Column names need to be UPPERCASE
- The word separator needs to be an underscore, such as in `COLUMN_NAME`

You can use the following dataset list and the `names()` function to get and set each of their column names, and convert them according to our defined naming convention.

```
[14]: dataset_list <- c('raw_bike_sharing_systems.csv', 'raw_seoul_bike_sharing.csv', 'raw_cities_weather_forecast.csv', 'raw_worldcities.csv')
```

TODO: Write a `for` loop to iterate over the above datasets and convert their column names

```
[15]: for (dataset_name in dataset_list){  
  # Read dataset  
  dataset <- read_csv(dataset_name)  
  # Standardized its columns:  
  
  # Convert all column names to uppercase  
  names(dataset) <- toupper(names(dataset))  
  # Replace any white space separators by underscores, using the str_replace_all function  
  names(dataset) <- str_replace_all(names(dataset), " ", "_")  
  # Save the dataset  
  write_csv(dataset, dataset_name, row.names=FALSE)  
}
```

Parsed with column specification:

```
cols(  
  COUNTRY = col_character(),  
  CITY = col_character(),  
  NAME = col_character(),  
  SYSTEM = col_character(),  
  OPERATOR = col_character(),  
  LAUNCH_YEAR = col_character()
```

```
  PRESSURE = col_double(),  
  HUMIDITY = col_double(),  
  WIND_SPEED = col_double(),  
  WIND_DEG = col_double(),  
  SEASON = col_character(),  
  FORECAST_DATETIME = col_datetime(format = "")  
)  
Parsed with column specification:  
cols(  
  CITY = col_character(),  
  CITY_ASCII = col_character(),  
  LAT = col_double(),  
  LNG = col_double(),  
  COUNTRY = col_character(),  
  ISO2 = col_character(),  
  ISO3 = col_character(),  
  ADMIN_NAME = col_character(),  
  CAPITAL = col_character(),  
  POPULATION = col_double(),  
  ID = col_double()  
)
```

TODO: Read the resulting datasets back and check whether their column names follow the naming convention

```
[16]: for (dataset_name in dataset_list){  
  # Print a summary for each data set to check whether the column names were correctly converted  
  dataset <- read_csv(dataset_name)  
  summary(dataset)  
}
```

Parsed with column specification:

```
cols(  
  COUNTRY = col_character(),  
  CITY = col_character(),  
  NAME = col_character(),  
  SYSTEM = col_character(),  
  OPERATOR = col_character(),  
  LAUNCHED = col_character(),  
  DISCONTINUED = col_character()
```


Data Wrangling (Regex)

```
#  
Parsed with column specification:  
cols(  
  CITY = col_character(),  
  CITY_ASCII = col_character(),  
  LAT = col_double(),  
  LNG = col_double(),  
  COUNTRY = col_character(),  
  ISO2 = col_character(),  
  ISO3 = col_character(),  
  ADMIN_NAME = col_character(),  
  CAPITAL = col_character(),  
  POPULATION = col_double(),  
  ID = col_double()  
)
```

Process the web-scraped bike sharing system dataset

By now we have standardized all column names. Next, we will focus on cleaning up the values in the web-scraped bike sharing systems dataset.

```
[17]: # First load the dataset  
bike_sharing_df <- read_csv("raw_bike_sharing_systems.csv")
```

```
Parsed with column specification:  
cols(  
  COUNTRY = col_character(),  
  CITY = col_character(),  
  NAME = col_character(),  
  SYSTEM = col_character(),  
  OPERATOR = col_character(),  
  LAUNCHED = col_character(),  
  DISCONTINUED = col_character(),  
  STATIONS = col_character(),  
  BICYCLES = col_character(),  
  DAILY_RIDERSHIP = col_character()  
)
```

```
[18]: # Print (to head)  
head(bike_sharing_df)
```

A tibble: 6 × 10									
COUNTRY	CITY	NAME	SYSTEM	OPERATOR	LAUNCHED	DISCONTINUED	STATIONS	BICYCLES	DAILY_RIDERSHIP
<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
Albania	Tirana	Ecovolis	NA	NA	March 2011	NA	8	200	NA
Argentina	Mendoza	Metrobici	NA	NA	2014	NA	2	40	NA
Argentina	San Lorenzo, Santa Fe	Bicludad	Bicludad	NA	27 November 2016	NA	8	80	NA
Argentina	Buenos Aires	Ecobici	Serttel Brasil	Bike in Baires Consortium,[10]	2010	NA	400	4000	21917
Argentina	Rosario	Mi Bici Tu Bici[11]	NA	NA	2 December 2015	NA	47	480	NA
Australia	Melbourne[12]	Melbourne Bike Share	PBSC & 8D	Motivate	June 2010	30 November 2019[13]	53	676	NA

Even from the first few rows, you can see there is plenty of undesirable embedded textual content, such as the reference link included in [Melbourne\[12\]](#).

In this project, let's only focus on processing the following relevant columns (feel free to process the other columns for more practice):

- **COUNTRY**: Country name
- **CITY**: City name
- **SYSTEM**: Bike-sharing system name
- **BICYCLES**: Total number of bikes in the system

```
[19]: # Select the four columns  
sub_bike_sharing_df <- bike_sharing_df %>% select(COUNTRY, CITY, SYSTEM, BICYCLES)
```

Data Wrangling (Regex)

```
Let's see the types of the selected columns
```

```
[20]: sub_bike_sharing_df %>%
      summarise_all(class) %>%
      gather(variable, class)
```

A tibble: 4 × 2

variable	class
COUNTRY	character
CITY	character
SYSTEM	character
BICYCLES	character

They are all interpreted as character columns, but we expect the `BICYCLES` column to be of numeric type. Let's see why it wasn't loaded as a numeric column - possibly some entries contain characters. Let's create a simple function called `find_character` to check that.

```
[21]: # grepl searches a string for non-digital characters, and returns TRUE or FALSE
      # if it finds any non-digital characters, then the bicycle column is not purely numeric
      find_character <- function(strings) grepl("[^0-9]", strings)
```

Let's try to find any elements in the `Bicycles` column containing non-numeric characters.

```
[22]: sub_bike_sharing_df %>%
      select(BICYCLES) %>%
      filter(find_character(BICYCLES)) %>%
      slice(0:10)
```

A spec_tbl_df: 10 × 1

```
Let's see the types of the selected columns
```

```
[20]: sub_bike_sharing_df %>%
      summarise_all(class) %>%
      gather(variable, class)
```

A tibble: 4 × 2

variable	class
COUNTRY	character
CITY	character
SYSTEM	character
BICYCLES	character

They are all interpreted as character columns, but we expect the `BICYCLES` column to be of numeric type. Let's see why it wasn't loaded as a numeric column - possibly some entries contain characters. Let's create a simple function called `find_character` to check that.

```
[21]: # grepl searches a string for non-digital characters, and returns TRUE or FALSE
      # if it finds any non-digital characters, then the bicycle column is not purely numeric
      find_character <- function(strings) grepl("[^0-9]", strings)
```

Let's try to find any elements in the `Bicycles` column containing non-numeric characters.

```
[22]: sub_bike_sharing_df %>%
      select(BICYCLES) %>%
      filter(find_character(BICYCLES)) %>%
      slice(0:10)
```

A spec_tbl_df: 10 × 1

Data Wrangling (Regex)

```
filter(find_character(BICYCLES)) %>%
  slice(0:10)
```

A spec_tbl_df: 10 x 1

BICYCLES
<chr>
4115[22]
310[59]
500[72]
[75]
180[76]
600[77]
[78]
Initially 800 (later 2500)
100 (220)
370[144]

As you can see, many rows have non-numeric characters, such as `32` (including 6 rollers) `[162]` and `1000[253]`. This is actually very common for a table scraped from Wiki when no input validation is enforced.

Next, let's take a look at the other columns, namely `COUNTRY`, `CITY`, and `SYSTEM`, to see if they contain any undesired reference links, such as in `Melbourne[12]`.

```
[23]: # Define a 'reference link' character class,
# '[A-z0-9]' means at least one character
# '[' and '\]' means the character is wrapped by [], such as for [12] or [abc]
ref_pattern <- "\\[[A-z0-9]+\\]"
find_reference_pattern <- function(strings) grepl(ref_pattern, strings)
```

```
[24]: # Check whether the COUNTRY column has any reference links
sub_bike_sharing_df %>%
  select(COUNTRY) %>%
  filter(find_reference_pattern(COUNTRY)) %>%
  slice(0:10)
```

A
spec_tbl_df:
0 x 1

COUNTRY

<chr>

Ok, looks like the `COUNTRY` column is clean. Let's check the `CITY` column.

```
[25]: # Check whether the CITY column has any reference links
sub_bike_sharing_df %>%
  select(CITY) %>%
  filter(find_reference_pattern(CITY)) %>%
  slice(0:10)
```

A spec_tbl_df: 10 x 1

CITY

<chr>

Melbourne[12]

Data Wrangling (Regex)

```
EasyBike[58]
4 Gen.[61]
3 Gen. SmoooveKey[113]
3 Gen. Smooove[141][142][143][139]
3 Gen. Smooove[179]
3 Gen. Smooove[181]
3 Gen. Smooove[183]
```

So the `SYSTEM` column also has some reference links.

After some preliminary investigations, we identified that the `CITY` and `SYSTEM` columns have some undesired reference links, and the `BICYCLES` column has both reference links and some textual annotations.

Next, you need to use regular expressions to clean up the unexpected reference links and text annotations in numeric values.

TASK: Remove undesired reference links using regular expressions

TODO: Write a custom function using `stringr::str_replace_all` to replace all reference links with an empty character for columns `CITY` and `SYSTEM`

```
[27]: # remove reference link
remove_ref <- function(strings) {
  ref_pattern <- "\\[[A-Z0-9]+\\]"
  # Replace all matched substrings with a white space using str_replace_all()
  strings <- str_replace_all(strings, ref_pattern, " ")
  return(strings)
}
```

TODO: Use the `dplyr::mutate()` function to apply the `remove_ref` function to the `CITY` and `SYSTEM` columns

```
result <- sub_bike_sharing_df %>% mutate(CITY=remove_ref(CITY),
                                         SYSTEM=remove_ref(SYSTEM),
                                         BICYCLES=remove_ref(BICYCLES))
```

TODO: Use the following code to check whether all reference links are removed:

```
result %>%
  select(CITY, SYSTEM, BICYCLES) %>%
  filter(find_reference_pattern(CITY) | find_reference_pattern(SYSTEM) | find_reference_pattern(BICYCLES))
```

A spec_tbl_df: 0 x 3

CITY SYSTEM BICYCLES

<chr> <chr> <chr>

Data Wrangling (Regex)

TASK: Extract the numeric value using regular expressions

TODO: Write a custom function using `stringr::str_extract` to extract the first digital substring match and convert it into numeric type For example, extract the value '32' from '32 (including 6 rollers)' [162].

```
[30]: # Extract the first number
extract_num <- function(columns){
  # Define a digital pattern
  digital_pattern <- "[0-9]+"
  # Find the first match using str_extract
  columns <- str_extract(columns, digital_pattern)
  # Convert the result to numeric using the as.numeric() function
  columns <- as.numeric(columns)
}
```

TODO: Use the `dplyr::mutate()` function to apply `extract_num` on the `BICYCLES` column

```
[31]: # Use the mutate() function on the BICYCLES column
result <- result %>%
  mutate(BICYCLES = extract_num(BICYCLES))
```

TODO: Use the summary function to check the descriptive statistics of the numeric `BICYCLES` column

```
[32]: summary(result$BICYCLES)
      Min. 1st Qu.  Median    Mean 3rd Qu.     Max.   NA's
      5      100      350    2022   1400    78000      78
```

TODO: Write the cleaned bike-sharing systems dataset into a csv file called `bike_sharing_systems.csv`

```
[34]: # Write dataset to 'bike_sharing_systems.csv'
write.csv(result, file = "bike_sharing_systems.csv", row.names = FALSE)
```

EDA With SQL

Establish your SQLite connection

Load the 'RSQLite' library, and use the 'dbConnect()' function as you did in the previous lab to establish the connection to your SQLite database.

You are now ready to start running SQL queries using the RSQLite library as you did in Course 3.

```
[2]: # provide your solution here
con <- dbConnect(SQLite(), "seoul.db")
```

```
[4]: library(readr)
```

```
Warning message:
"replacing previous import 'lifecycle::last_warnings' by 'rlang::last_warnings' when loading 'tibble'"
Warning message:
"replacing previous import 'ellipsis::check_dots_unnamed' by 'rlang::check_dots_unnamed' when loading 'tibble'"
Warning message:
"replacing previous import 'ellipsis::check_dots_used' by 'rlang::check_dots_used' when loading 'tibble'"
Warning message:
"replacing previous import 'ellipsis::check_dots_empty' by 'rlang::check_dots_empty' when loading 'tibble'"
```

```
[5]: # Download the CSV files
world_cities <- read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-RP0321EN-SkillsNetwork/
bike_sharing_systems <- read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-RP0321EN-SkillsNetwork/
cities_weather_forecast <- read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-RP0321EN-SkillsNetwork/
seoul_bike_sharing <- read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-RP0321EN-SkillsNetwork/

# Write the data to the SQLite database
dbWriteTable(con, "WORLD_CITIES", world_cities, overwrite = TRUE)
dbWriteTable(con, "BIKE_SHARING_SYSTEMS", bike_sharing_systems, overwrite = TRUE)
dbWriteTable(con, "CITIES_WEATHER_FORECAST", cities_weather_forecast, overwrite = TRUE)
dbWriteTable(con, "SEOUL_BIKE_SHARING", seoul_bike_sharing, overwrite = TRUE)
```

```
Parsed with column specification:
cols(
  CITY = col_character(),
  CITY_ASCII = col_character(),
```

```
[5]: # Execute SQL queries to verify data loading
result <- dbGetQuery(con, "SELECT * FROM WORLD_CITIES LIMIT 5")
print("WORLD_CITIES table:")
print(result)
```

```
[1] "WORLD_CITIES table:"
  CITY CITY_ASCII  LAT  LNG  COUNTRY ISO2 ISO3 ADMIN_NAME CAPITAL
1 Tokyo Tokyo 35.6897 139.6922 Japan JP JPN Tōkyō primary
2 Jakarta Jakarta -6.2146 106.8451 Indonesia ID IDN Jakarta primary
3 Delhi Delhi 28.6600 77.2300 India IN IND Delhi admin
4 Mumbai Mumbai 18.9667 72.8333 India IN IND Mahārāshtra admin
5 Manila Manila 14.5958 120.9772 Philippines PH PHL Manila primary
  POPULATION ID
1 37977000 1392685764
2 34540000 1360771077
3 29617000 1356872604
4 23355000 1356226629
5 23088000 1608618140
```

EDA With SQL

Task 1 - Record Count

Determine how many records are in the seoul_bike_sharing dataset.

Solution 1

```
9]: # provide your solution here
query <- "
SELECT DATE, HOUR, SUM(RENTED_BIKE_COUNT) AS total_rentals
FROM seoul_bike_sharing
GROUP BY DATE, HOUR
ORDER BY total_rentals DESC
LIMIT 1;
"
result <- dbGetQuery(con, query)
result
```

A data.frame: 1 × 3

DATE	HOUR	total_rentals
<chr>	<dbl>	<dbl>
19/06/2018	18	3556

Task 7 - Hourly popularity and temperature by season

Determine the average hourly temperature and the average number of bike rentals per hour over each season. List the top ten results by average bike count.

Solution 7

```
[8]: query <- "
SELECT s.SEASONS,
      AVG(s.TEMPERATURE) AS avg_hourly_temperature,
      AVG(s.RENTED_BIKE_COUNT) AS avg_bike_rentals_per_hour
FROM seoul_bike_sharing s
GROUP BY s.SEASONS
ORDER BY avg_bike_rentals_per_hour DESC
LIMIT 10;
"
result <- dbGetQuery(con, query)
result
```

A data.frame: 4 × 3

SEASONS	avg_hourly_temperature	avg_bike_rentals_per_hour
<chr>	<dbl>	<dbl>
Summer	26.587711	1034.0734
Autumn	13.821580	924.1105
Spring	13.021685	746.2542
Winter	-2.540463	225.5412

EDA With SQL

Find the average hourly bike count during each season.

Also include the minimum, maximum, and standard deviation of the hourly bike count for each season.

Hint : Use the $\text{SQRT}(\text{AVG}(\text{col}^2) - \text{AVG}(\text{col})^2)$ function where col refers to your column name for finding the standard deviation

Solution 8

[12]: # provide your solution here

```
query<- "
SELECT
  SEASONS,
  AVG(RENTED_BIKE_COUNT) AS avg_hourly_bike_count,
  MIN(RENTED_BIKE_COUNT) AS min_hourly_bike_count,
  MAX(RENTED_BIKE_COUNT) AS max_hourly_bike_count,
  SQRT(AVG(RENTED_BIKE_COUNT * RENTED_BIKE_COUNT) - AVG(RENTED_BIKE_COUNT) * AVG(RENTED_BIKE_COUNT)) AS std_dev_hourly_bike_count
FROM
  seoul_bike_sharing
GROUP BY
  SEASONS;
"
result <- dbGetQuery(con, query)
print(result)
```

	SEASONS	avg_hourly_bike_count	min_hourly_bike_count	max_hourly_bike_count
1	Autumn	924.1105	2	3298
2	Spring	746.2542	2	3251
3	Summer	1034.0734	9	3556
4	Winter	225.5412	3	937

	std_dev_hourly_bike_count
1	617.3885
2	618.5247
3	690.0884

[11]: # Weather Seasonality

```
query <- "
SELECT
  SEASONS,
  AVG(TEMPERATURE) AS avg_temperature,
  AVG(HUMIDITY) AS avg_humidity,
  AVG(WIND_SPEED) AS avg_wind_speed,
  AVG(VISIBILITY) AS avg_visibility,
  AVG(DEW_POINT_TEMPERATURE) AS avg_dew_point_temperature,
  AVG(SOLAR_RADIATION) AS avg_solar_radiation,
  AVG(RAINFALL) AS avg_rainfall,
  AVG(SNOWFALL) AS avg_snowfall,
  AVG(RENTED_BIKE_COUNT) AS avg_bike_count
FROM
  seoul_bike_sharing
GROUP BY
  SEASONS
ORDER BY
  avg_bike_count DESC;
"
```

```
result <- dbGetQuery(con, query)
print(result)
```

	SEASONS	avg_temperature	avg_humidity	avg_wind_speed	avg_visibility
1	Summer	26.587711	64.98143	1.609420	1501.745
2	Autumn	13.821580	59.04491	1.492101	1558.174
3	Spring	13.021685	58.75833	1.857778	1240.912
4	Winter	-2.540463	49.74491	1.922685	1445.987

	avg_dew_point_temperature	avg_solar_radiation	avg_rainfall	avg_snowfall
1	18.750136	0.7612545	0.25348732	0.00000000
2	5.150594	0.5227827	0.11765617	0.06350026
3	4.091389	0.6803009	0.18694444	0.00000000
4	-12.416667	0.2981806	0.03282407	0.24750000

	avg_bike_count
1	1034.0734
2	924.1105
3	746.2542
4	225.5412

EDA With SQL

Task 10 - Total Bike Count and City Info for Seoul

Use an implicit join across the `WORLD_CITIES` and the `BIKE_SHARING_SYSTEMS` tables to determine the total number of bikes available in Seoul, plus the following city information about Seoul: `CITY`, `COUNTRY`, `LAT`, `LO`, `POPULATION`, in a single view.

Notice that in this case, the `CITY` column will work for the `WORLD_CITIES` table, but in general you would have to use the `CITY_ASCII` column.

Solution 10

```
[8]: # Execute SQL query to retrieve city information and total bikes available in Seoul
query <- "
SELECT
  wc.CITY,
  wc.COUNTRY,
  wc.LAT,
  wc.LNG AS LO,
  wc.POPULATION,
  bs.BICYCLES AS total_bikes_available
FROM
  WORLD_CITIES wc
JOIN
  BIKE_SHARING_SYSTEMS bs ON wc.CITY_ASCII = bs.CITY
WHERE
  wc.CITY = 'Seoul';
"
result <- dbGetQuery(con, query)
print(result)

  CITY      COUNTRY      LAT LO POPULATION total_bikes_available
1 Seoul Korea, South 37.5833 127  21794000          20000
```

Solution 11

```
•[10]: # provide your solution here
# Execute SQL query to retrieve city names and coordinates with total bike counts between 15000 and 20000
query <- "
SELECT
  wc.CITY,
  wc.COUNTRY,
  wc.LAT,
  wc.LNG AS LNG,
  wc.POPULATION,
  bs.BICYCLES AS total_bikes
FROM
  WORLD_CITIES wc
JOIN
  BIKE_SHARING_SYSTEMS bs ON wc.CITY_ASCII = bs.CITY
WHERE
  bs.BICYCLES BETWEEN 15000 AND 20000;
"

# Execute the query
result <- dbGetQuery(con, query)

# Print the result
print(result)
```

	CITY	COUNTRY	LAT	LNG	POPULATION	total_bikes
1	Beijing	China	39.9050	116.3914	19439000	16000
2	Ningbo	China	29.8750	121.5492	7639000	15000
3	Shanghai	China	31.1667	121.4667	22120000	19165
4	Weifang	China	36.7167	119.1000	9373000	20000
5	Xi'an	China	34.2667	108.9000	7135000	20000
6	Zhuzhou	China	27.8407	113.1469	3855609	20000
7	Seoul Korea, South		37.5833	127.0000	21794000	20000

```
[ ]: close(conn)
```

EDA With Data Visualization

Task 1 - Load the dataset

Ensure you read `DATE` as type `character`.

Solution 1

```
[4]: seoul_bike_sharing_url <- "https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBMDeveloperSkillsNetwork-RP0321EN-SkillsNetwork,

# Load the dataset with DATE as character type
seoul_bike_sharing <- read.csv(seoul_bike_sharing_url, colClasses = c(DATE = "character"))

# Display the structure of the dataset
str(seoul_bike_sharing)
```

```
'data.frame': 8465 obs. of 14 variables:
 $ DATE      : chr "01/12/2017" "01/12/2017" "01/12/2017" "01/12/2017" ...
 $ RENTED_BIKE_COUNT : int 254 204 173 107 78 100 181 460 930 490 ...
 $ HOUR       : int 0 1 2 3 4 5 6 7 8 9 ...
 $ TEMPERATURE : num -5.2 -5.5 -6 -6.2 -6 -6.4 -6.6 -7.4 -7.6 -6.5 ...
 $ HUMIDITY    : int 37 38 39 40 36 37 35 38 37 27 ...
 $ WIND_SPEED  : num 2.2 0.8 1 0.9 2.3 1.5 1.3 0.9 1.1 0.5 ...
 $ VISIBILITY  : int 2000 2000 2000 2000 2000 2000 2000 2000 2000 1928 ...
 $ DEW_POINT_TEMPERATURE: num -17.6 -17.6 -17.7 -17.6 -18.6 -18.7 -19.5 -19.3 -19.8 -22.4 ...
 $ SOLAR_RADIATION : num 0 0 0 0 0 0 0 0.01 0.23 ...
 $ RAINFALL    : num 0 0 0 0 0 0 0 0 0 ...
 $ SNOWFALL    : num 0 0 0 0 0 0 0 0 0 ...
 $ SEASONS     : Factor w/ 4 levels "Autumn","Spring",...: 4 4 4 4 4 4 4 4 4 ...
 $ HOLIDAY     : Factor w/ 2 levels "Holiday","No Holiday": 2 2 2 2 2 2 2 2 2 ...
 $ FUNCTIONING_DAY : Factor w/ 1 level "Yes": 1 1 1 1 1 1 1 1 1 ...
```

Task 2 - Recast `DATE` as a date

Use the format of the data, namely `"%d/%m/%Y"`.

Solution 2

```
[5]: # Recast DATE as a date with the format "%d/%m/%Y"
seoul_bike_sharing$DATE <- as.Date(seoul_bike_sharing$DATE, format = "%d/%m/%Y")

# Display the structure of the updated dataset
str(seoul_bike_sharing)
```

```
'data.frame': 8465 obs. of 14 variables:
 $ DATE      : Date, format: "2017-12-01" "2017-12-01" ...
 $ RENTED_BIKE_COUNT : int 254 204 173 107 78 100 181 460 930 490 ...
 $ HOUR       : int 0 1 2 3 4 5 6 7 8 9 ...
 $ TEMPERATURE : num -5.2 -5.5 -6 -6.2 -6 -6.4 -6.6 -7.4 -7.6 -6.5 ...
 $ HUMIDITY    : int 37 38 39 40 36 37 35 38 37 27 ...
 $ WIND_SPEED  : num 2.2 0.8 1 0.9 2.3 1.5 1.3 0.9 1.1 0.5 ...
 $ VISIBILITY  : int 2000 2000 2000 2000 2000 2000 2000 2000 2000 1928 ...
 $ DEW_POINT_TEMPERATURE: num -17.6 -17.6 -17.7 -17.6 -18.6 -18.7 -19.5 -19.3 -19.8 -22.4 ...
 $ SOLAR_RADIATION : num 0 0 0 0 0 0 0 0.01 0.23 ...
 $ RAINFALL    : num 0 0 0 0 0 0 0 0 0 ...
 $ SNOWFALL    : num 0 0 0 0 0 0 0 0 0 ...
 $ SEASONS     : Factor w/ 4 levels "Autumn","Spring",...: 4 4 4 4 4 4 4 4 4 ...
 $ HOLIDAY     : Factor w/ 2 levels "Holiday","No Holiday": 2 2 2 2 2 2 2 2 2 ...
 $ FUNCTIONING_DAY : Factor w/ 1 level "Yes": 1 1 1 1 1 1 1 1 1 ...
```

EDA With Data Visualization

Task 3 - Cast `HOURS` as a categorical variable

Also, coerce its levels to be an ordered sequence. This will ensure your visualizations correctly utilize `HOURS` as a discrete variable with the expected ordering.

Solution 3

```
[6]: # provide your solution here
# Cast HOUR as a categorical variable with ordered levels
seoul_bike_sharing$HOUR <- factor(seoul_bike_sharing$HOUR, levels = 0:23, ordered = TRUE)
```

Check the structure of the dataframe

```
[7]: str(seoul_bike_sharing)

'data.frame':  8465 obs. of  14 variables:
 $ DATE          : Date, format: "2017-12-01" "2017-12-01" ...
 $ RENTED_BIKE_COUNT : int  254 204 173 107 78 100 181 460 930 490 ...
 $ HOUR          : Ord.factor w/ 24 levels "0"<"1"<"2"<"3"<...: 1 2 3 4 5 6 7 8 9 10 ...
 $ TEMPERATURE    : num  -5.2 -5.5 -6 -6.2 -6 -6.4 -6.6 -7.4 -7.6 -6.5 ...
 $ HUMIDITY        : int   37 38 39 40 36 37 35 38 37 27 ...
 $ WIND_SPEED      : num   2.2 0.8 1 0.9 2.3 1.5 1.3 0.9 1.1 0.5 ...
 $ VISIBILITY      : int  2000 2000 2000 2000 2000 2000 2000 2000 1928 ...
 $ DEW_POINT_TEMPERATURE: num  -17.6 -17.6 -17.7 -17.6 -18.6 -18.7 -19.5 -19.3 -19.8 -22.4 ...
 $ SOLAR_RADIATION : num   0 0 0 0 0 0 0 0.01 0.23 ...
 $ RAINFALL        : num   0 0 0 0 0 0 0 0 ...
 $ SNOWFALL        : num   0 0 0 0 0 0 0 0 ...
 $ SEASONS         : Factor w/ 4 levels "Autumn","Spring",...: 4 4 4 4 4 4 4 4 4 ...
 $ HOLIDAY         : Factor w/ 2 levels "Holiday","No Holiday": 2 2 2 2 2 2 2 2 2 ...
 $ FUNCTIONING_DAY : Factor w/ 1 level "Yes": 1 1 1 1 1 1 1 1 1 ...
```

Finally, ensure there are no missing values

```
[8]: sum(is.na(seoul_bike_sharing))

0
```

Descriptive Statistics

Now you are all set to take a look at some high level statistics of the `seoul_bike_sharing` dataset.

Task 4 - Dataset Summary

Use the base R `summary()` function to describe the `seoul_bike_sharing` dataset.

Solution 4

```
[9]: # provide your solution here
summary(seoul_bike_sharing)
```

DATE	RENTED_BIKE_COUNT	HOUR	TEMPERATURE
Min. :2017-12-01	Min. : 2.0	7	Min. : -17.80
1st Qu.:2018-02-27	1st Qu.: 214.0	8	1st Qu.: 3.00
Median :2018-05-28	Median : 542.0	9	Median : 13.50
Mean :2018-05-28	Mean : 729.2	10	Mean : 12.77
3rd Qu.:2018-08-24	3rd Qu.:1084.0	11	3rd Qu.: 22.70
Max. :2018-11-30	Max. :3556.0	12	Max. : 39.40
		(Other):6347	
HUMIDITY	WIND_SPEED	VISIBILITY	DEW_POINT_TEMPERATURE
Min. : 0.00	Min. :0.000	Min. : 27	Min. : -30.600
1st Qu.:42.00	1st Qu.:0.900	1st Qu.: 935	1st Qu.: -5.100
Median :57.00	Median :1.500	Median :1690	Median : 4.700
Mean :58.15	Mean :1.726	Mean :1434	Mean : 3.945
3rd Qu.:74.00	3rd Qu.:2.300	3rd Qu.:2000	3rd Qu.: 15.200
Max. :98.00	Max. :7.400	Max. :2000	Max. : 27.200

EDA With Data Visualization

```
PREVAILING : 1.7 m/s    PRECIPITATION : 1.726    PRECIPITATION : 1434    PRECIPITATION : 3.945
Mean :58.15    Mean :1.726    Mean :1434    Mean : 3.945
3rd Qu.:74.00    3rd Qu.:2.300    3rd Qu.:2000    3rd Qu.: 15.200
Max. :98.00    Max. :7.400    Max. :2000    Max. : 27.200

SOLAR_RADIATION    RAINFALL    SNOWFALL    SEASONS
Min. :0.0000    Min. : 0.0000    Min. :0.00000    Autumn:1937
1st Qu.:0.0000    1st Qu.: 0.0000    1st Qu.:0.00000    Spring:2160
Median :0.0100    Median : 0.0000    Median :0.00000    Summer:2208
Mean :0.5679    Mean : 0.1491    Mean :0.07769    Winter:2160
3rd Qu.:0.9300    3rd Qu.: 0.0000    3rd Qu.:0.00000
Max. :3.5200    Max. :35.0000    Max. :8.80000

HOLIDAY    FUNCTIONING_DAY
Holiday : 408    Yes:8465
No Holiday:8057
```

Some Basic Observations:

- We can see from `DATE` that we have exactly a full year of data.
- No records have zero bike counts.
- Spring and Winter have the same count of records, while autumn has the least and Summer has the most.
- Temperature has a large range, so we might expect it to explain at least some of the variation in bike rentals.
- Precipitation seems to be quite rare, only happening in the fourth quartiles for both `RAINFALL` and `SNOWFALL`.
- The average `WINDSPEED` is very light at only 1.7 m/s, and even the maximum is only a moderate breeze (Google 'Beaufort Wind Scale' to find the different wind descriptions)

By now, you might agree that Exploratory Data Analysis can create more questions than answers. That's okay - you'll have a much deeper understanding and appreciation for your data as a result!

Task 5 - Based on the above stats, calculate how many Holidays there are.

Solution 5:

```
[10]: # provide your solution here

# Calculate the number of holidays
holiday_count <- table(seoul_bike_sharing$HOLIDAY)["Holiday"]

# Display the result
holiday_count
```

Holiday: 408

Task 6 - Calculate the percentage of records that fall on a holiday.

Solution 6

```
[11]: # provide your solution here

# Calculate the total number of records
total_records <- nrow(seoul_bike_sharing)

# Calculate the percentage
percentage_holidays <- (holiday_count / total_records) * 100

# Display the result
percentage_holidays
```

Holiday: 4.8198464264619

EDA With Data Visualization

Task 7 - Given there is exactly a full year of data, determine how many records we expect to have.

Solution 7

```
[13]: # provide your solution here

# Assuming hourly data and a full year
days_in_year <- 365
hours_per_day <- 24

# Calculate expected number of records
expected_records <- days_in_year * hours_per_day

# Print the result
cat("Expected Number of Records:", expected_records, "\n")
```

Expected Number of Records: 8760

Task 8 - Given the observations for the 'FUNCTIONING_DAY' how many records must there be?

Solution 8

```
[15]: # provide your solution here


# Count the number of records for each level in FUNCTIONING_DAY
table(seoul_bike_sharing$FUNCTIONING_DAY)

Yes
8465
```

```
[20]: # SCATTER PLOT OF RENTED_BIKE_COUNT VS DATE
```

```
ggplot(seoul_bike_sharing, aes(x = DATE, y = RENTED_BIKE_COUNT)) +
  geom_point(alpha = 0.5) + # Set opacity using alpha parameter
  labs(title = "Scatter Plot of RENTED_BIKE_COUNT vs DATE",
       x = "Date",
       y = "Rented Bike Count") +
  theme_minimal()
```

Scatter Plot of RENTED_BIKE_COUNT vs DATE




Ungraded Task: We can see some patterns emerging here.

EDA With Data Visualization

[21]: # provide your solution here

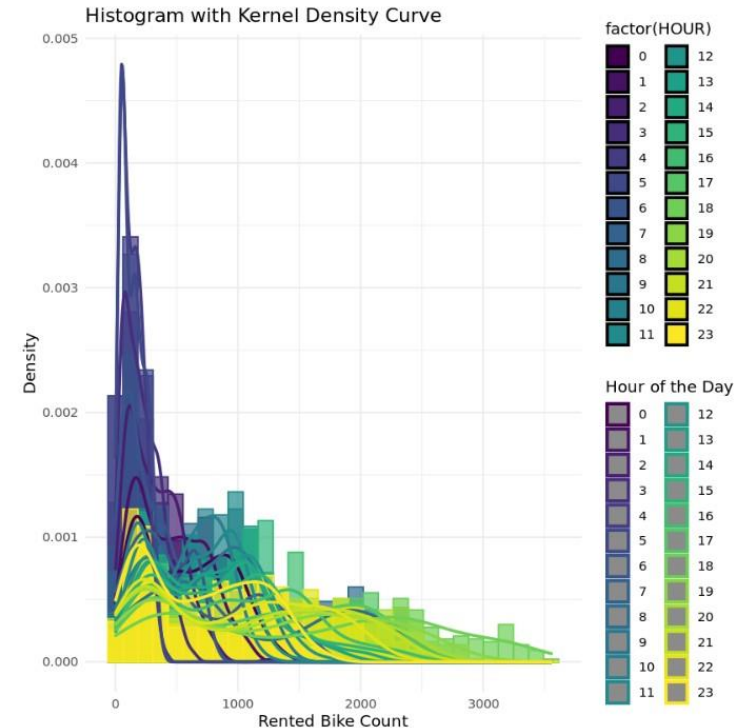
```
ggplot(seoul_bike_sharing, aes(x = DATE, y = RENTED_BIKE_COUNT, color = factor(HOUR))) +  
  geom_point(alpha = 0.5) + # Set opacity using alpha parameter  
  labs(title = "Scatter Plot of RENTED_BIKE_COUNT vs DATE with Color by HOURS",  
       x = "Date",  
       y = "Rented Bike Count",  
       color = "Hour of the Day") +  
  theme_minimal()
```

Scatter Plot of RENTED_BIKE_COUNT vs DATE with Color by HOURS



[22]: # provide your solution here

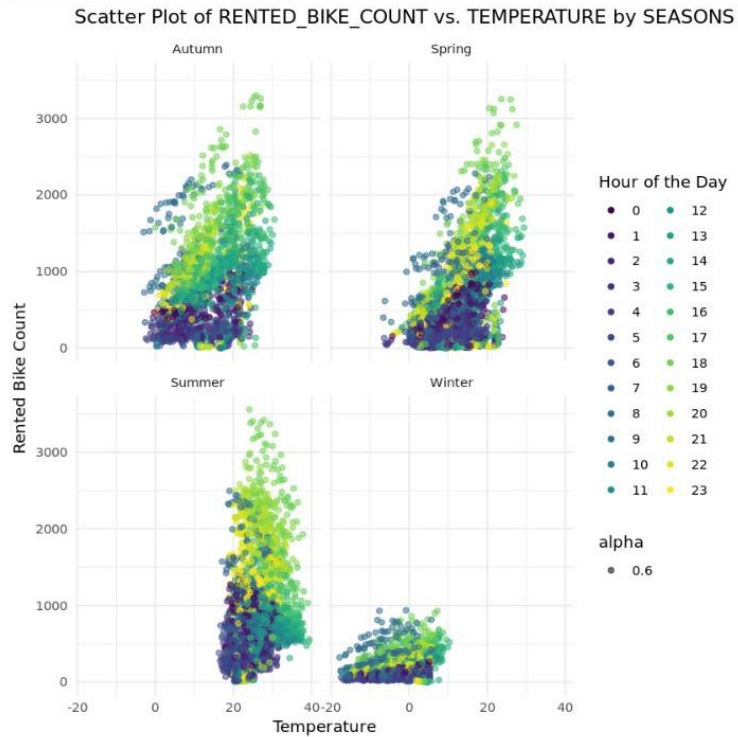
```
ggplot(seoul_bike_sharing, aes(x = RENTED_BIKE_COUNT, y = ..density..)) +  
  geom_histogram(aes(fill = factor(HOUR), color = factor(HOUR)),  
               bins = 30, alpha = 0.7, position = "identity") +  
  geom_density(aes(color = factor(HOUR)), size = 1) +  
  labs(title = "Histogram with Kernel Density Curve",  
       x = "Rented Bike Count",  
       y = "Density",  
       color = "Hour of the Day") +  
  theme_minimal()
```



EDA With Data Visualization


[23]: # provide your solution here

```
ggplot(seoul_bike_sharing, aes(x = TEMPERATURE, y = RENTED_BIKE_COUNT, color = factor(HOUR), alpha = 0.6)) +  
  geom_point() +  
  facet_wrap(~SEASONS) +  
  labs(title = "Scatter Plot of RENTED_BIKE_COUNT vs. TEMPERATURE by SEASONS",  
       x = "Temperature",  
       y = "Rented Bike Count",  
       color = "Hour of the Day") +  
  theme_minimal()
```



Comparing this plot to the same plot below, but without grouping by SEASONS, shows how important seasonality is in explaining bike rental counts.

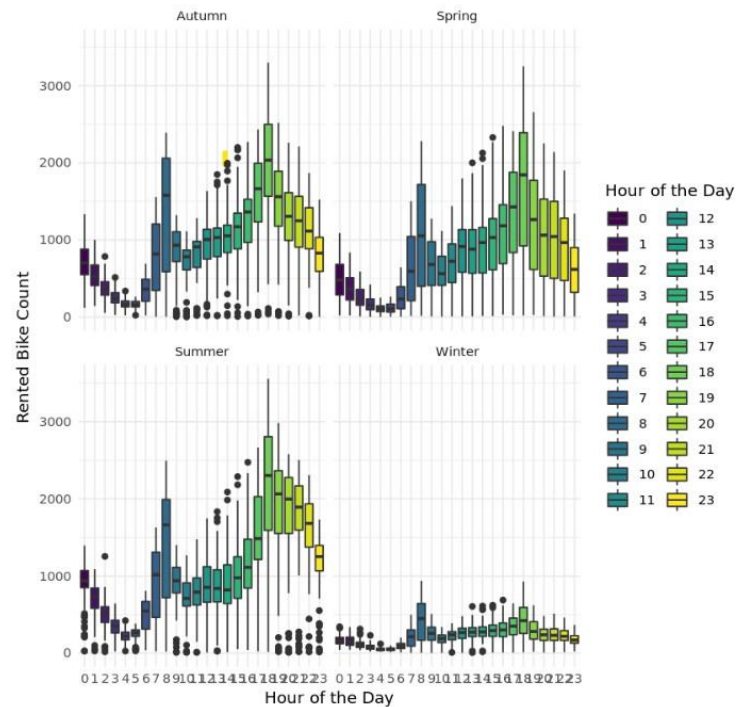
[24]: ggplot(seoul_bike_sharing) +
 geom_point(aes(x=TEMPERATURE,y=RENTED_BIKE_COUNT,colour=HOUR),alpha=1/5)



EDA With Data Visualization

```
[26]: # provide your solution here
ggplot(seoul_bike_sharing, aes(x = factor(HOUR), y = RENTED_BIKE_COUNT, fill = factor(HOUR))) +
  geom_boxplot() +
  facet_wrap(~SEASONS) +
  labs(title = "Boxplots of RENTED_BIKE_COUNT vs. HOUR by SEASONS",
       x = "Hour of the Day",
       y = "Rented Bike Count",
       fill = "Hour of the Day") +
  theme_minimal()
```

Boxplots of RENTED_BIKE_COUNT vs. HOUR by SEASONS



Task 15 - Group the data by DATE, and use the summarize() function to calculate the daily total rainfall and snowfall.

Also, go ahead and plot the results if you wish.

Solution 15

```
[27]: # provide your solution here
daily_summary <- seoul_bike_sharing %>%
  group_by(DATE) %>%
  summarize(total_rainfall = sum(RAINFALL),
            total_snowfall = sum(SNOWFALL))

# Plot the results
ggplot(daily_summary, aes(x = DATE)) +
  geom_line(aes(y = total_rainfall, color = "Total Rainfall"), size = 1.5) +
  geom_line(aes(y = total_snowfall, color = "Total Snowfall"), size = 1.5) +
  labs(title = "Daily Total Rainfall and Snowfall",
       x = "Date",
       y = "Total Amount",
       color = "Type") +
  scale_color_manual(values = c("Total Rainfall" = "blue", "Total Snowfall" = "red")) +
  theme_minimal()
```

Daily Total Rainfall and Snowfall

